



P95 TTFT vs load: three serving configurations

Median per-session P95 TTFT (ms)


● Marconi (FLOP-aware eviction) ◆ LRU prefix cache ■ No prefix cache



(a) LMSys



(b) ShareGPT



(c) SWEBench

Session arrival rate (sessions/s) — cf. Marconi Fig. 13(a)