




Marconi vs LRU (both with prefix cache on)
Ratio = P95 TTFT Marconi / P95 TTFT LRU per session; <1 means Marconi faster



(a) LMSys



(b) ShareGPT



(c) SWEBench