


SWEBench: three configurations vs load (inter-request time $\propto$ art)


Median per-session P95 TTFT (ms)

● Marconi (FLOP-aware eviction) ◆ LRU prefix cache ■ No prefix cache

art = 5.0 s



art = 7.5 s



art = 10.0 s

